


Tokens spent per problem, 200-problem subset x 5 models (pooled, n = 1,000 per arm)

Non-thinking (generation estimated)



Thinking (exact)

